

HousingProj

September 12, 2021

```
[1]: #!/pylab inline
import pandas as pd
import numpy as np

df = pd.read_csv('Housing Project/housingproj.csv')
#print(df.head(5))
print(df.shape)

#FILL the missing Values
df.isnull().sum()
```

(20640, 10)

```
[1]: longitude          0
latitude              0
housing_median_age    0
total_rooms           0
total_bedrooms       207
population            0
households            0
median_income         0
ocean_proximity       0
median_house_value    0
dtype: int64
```

Data Dictionary:

Features: 'longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms', 'population', 'households', 'median_income', 'ocean_proximity'

Label: median_house_value

There are 20640 rows and 10 features. Since the response variable is continuous, it can be considered as linear regression problem.

```
[2]: #df.describe().transpose()
      #impute missing values with mean
      #df['total_bedrooms'].head()
      df.total_bedrooms.fillna(df.total_bedrooms.mean(),inplace=True)
      df.total_bedrooms.head()
      #print(tbdr_mean)
```

```
[2]: 0      129.0
      1     1106.0
      2      190.0
      3      235.0
      4      280.0
      Name: total_bedrooms, dtype: float64
```

```
[3]: from sklearn.preprocessing import LabelEncoder

      # Scikit Learn
      lb_temp = LabelEncoder()

      df['ocean_Dist'] = lb_temp.fit_transform(df.ocean_proximity)

      #print(df['ocean_Dist'])
      # Coding the labels. Converting the categories into the numbers.
      print(df.head(10))
      df1=df.drop(columns=['ocean_proximity'])
      print(df1.head())
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	
5	-122.25	37.85	52	919	213.0	
6	-122.25	37.84	52	2535	489.0	
7	-122.25	37.84	52	3104	687.0	
8	-122.26	37.84	42	2555	665.0	
9	-122.25	37.84	52	3549	707.0	

	population	households	median_income	ocean_proximity	median_house_value	\
0	322	126	8.3252	NEAR BAY	452600	
1	2401	1138	8.3014	NEAR BAY	358500	
2	496	177	7.2574	NEAR BAY	352100	
3	558	219	5.6431	NEAR BAY	341300	
4	565	259	3.8462	NEAR BAY	342200	
5	413	193	4.0368	NEAR BAY	269700	
6	1094	514	3.6591	NEAR BAY	299200	

7	1157	647	3.1200	NEAR BAY	241400
8	1206	595	2.0804	NEAR BAY	226700
9	1551	714	3.6912	NEAR BAY	261100

	ocean_Dist
0	3
1	3
2	3
3	3
4	3
5	3
6	3
7	3
8	3
9	3

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	

	population	households	median_income	median_house_value	ocean_Dist
0	322	126	8.3252	452600	3
1	2401	1138	8.3014	358500	3
2	496	177	7.2574	352100	3
3	558	219	5.6431	341300	3
4	565	259	3.8462	342200	3

```
[4]: #create dummy for Ocean_Dist

#dummy = pd.get_dummies(df1['ocean_Dist'],prefix='ocean_dist',drop_first=True)

#display(dummy)

df1 = df1.join(pd.
↳get_dummies(df1['ocean_Dist'],prefix='ocean_dist',drop_first=True))
display(df1)
df1.shape
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	
...	
20635	-121.09	39.48	25	1665	374.0	

20636	-121.21	39.49	18	697	150.0
20637	-121.22	39.43	17	2254	485.0
20638	-121.32	39.43	18	1860	409.0
20639	-121.24	39.37	16	2785	616.0

	population	households	median_income	median_house_value	ocean_Dist	\
0	322	126	8.3252	452600	3	
1	2401	1138	8.3014	358500	3	
2	496	177	7.2574	352100	3	
3	558	219	5.6431	341300	3	
4	565	259	3.8462	342200	3	
...	
20635	845	330	1.5603	78100	1	
20636	356	114	2.5568	77100	1	
20637	1007	433	1.7000	92300	1	
20638	741	349	1.8672	84700	1	
20639	1387	530	2.3886	89400	1	

	ocean_dist_1	ocean_dist_2	ocean_dist_3	ocean_dist_4
0	0	0	1	0
1	0	0	1	0
2	0	0	1	0
3	0	0	1	0
4	0	0	1	0
...
20635	1	0	0	0
20636	1	0	0	0
20637	1	0	0	0
20638	1	0	0	0
20639	1	0	0	0

[20640 rows x 14 columns]

[4]: (20640, 14)

```
[5]: #splitting to X and Y
X = df1.drop(columns = ['median_house_value',])
Y = df1[['median_house_value']]
print(Y.head())
print(X.head())
list(X.columns)
list(Y.columns)
```

	median_house_value
0	452600
1	358500
2	352100

```

3          341300
4          342200
   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0    -122.23    37.88                41          880          129.0
1    -122.22    37.86                21         7099         1106.0
2    -122.24    37.85                52         1467          190.0
3    -122.25    37.85                52         1274          235.0
4    -122.25    37.85                52         1627          280.0

   population  households  median_income  ocean_Dist  ocean_dist_1  \
0          322          126          8.3252          3          0
1         2401         1138          8.3014          3          0
2          496          177          7.2574          3          0
3          558          219          5.6431          3          0
4          565          259          3.8462          3          0

   ocean_dist_2  ocean_dist_3  ocean_dist_4
0              0              1              0
1              0              1              0
2              0              1              0
3              0              1              0
4              0              1              0

```

```
[5]: ['median_house_value']
```


```
[7]:
```

```

[6]: import seaborn as sns

#allow plot to appear inline
%matplotlib inline

sns.pairplot(df1, x_vars = ['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms',
                           'population', 'households', 'median_income'],
             y_vars='median_house_value', height=7, aspect =0.7, kind='scatter');
```



```
[7]: #Model Building
      #print(X.shape)
      #print(Y.shape)

      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler
      ss=StandardScaler()
      X_train,X_test, Y_train,Y_test = train_test_split(X,Y, test_size=0.3,
      ↪random_state=42)

      #X_standard = ss.
      ↪fit_transform(X_train[['longitude','latitude','housing_median_age','total_rooms','total_bed
      #
      ↪'population','households','median_income','ocean_dist_1','ocean_dist_2','ocean_dist_3','oce
      #display(X_standard[:10])
```

```
[8]: #display (X_standard[:10])
      #print ()
      display (X_test[:100])
      print ()
      display (Y_train[:10])
      print ()
      display (Y_test[:10])
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
20046	-119.01	36.06	25	1505	537.870553	
3024	-119.46	35.14	30	2943	537.870553	
15663	-122.44	37.80	52	3830	537.870553	
20484	-118.72	34.28	17	3051	537.870553	
9814	-121.93	36.62	34	2351	537.870553	
...
6052	-117.76	34.04	34	1914	537.870553	
15975	-122.45	37.77	52	2602	537.870553	
14331	-117.15	32.72	51	1321	537.870553	
1606	-122.08	37.88	26	2947	537.870553	
10915	-117.87	33.73	45	2264	537.870553	

	population	households	median_income	ocean_Dist	ocean_dist_1	\
20046	1392	359	1.6812	1	1	
3024	1565	584	2.5313	1	1	
15663	1310	963	3.4801	3	0	
20484	1705	495	5.7376	0	0	
9814	1063	428	3.7250	4	0	
...
6052	1564	328	2.8347	1	1	
15975	1330	647	3.5435	3	0	
14331	781	499	1.3071	4	0	

1606	825	626	2.9330	3	0
10915	1970	499	3.4193	0	0

	ocean_dist_2	ocean_dist_3	ocean_dist_4
20046	0	0	0
3024	0	0	0
15663	0	1	0
20484	0	0	0
9814	0	0	1
...
6052	0	0	0
15975	0	1	0
14331	0	0	1
1606	0	1	0
10915	0	0	0

[100 rows x 13 columns]

	median_house_value
7061	193800
14689	169700
17323	259800
10056	136100
15750	500001
15283	281300
18551	216100
7555	117400
15700	500001
12364	104200

	median_house_value
20046	47700
3024	45800
15663	500001
20484	218600
9814	278000
13311	158700
7113	198200
7668	157500
18246	340000
5723	446600

```
[9]: # split is 70% for training and 30% for testing
print(X_train.shape)
print(Y_train.shape)
print(X_test.shape)
print(Y_test.shape)
```

```
(14448, 13)
(14448, 1)
(6192, 13)
(6192, 1)
```

```
[5]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	
5	-122.25	37.85	52	919	213.0	
6	-122.25	37.84	52	2535	489.0	
7	-122.25	37.84	52	3104	687.0	
8	-122.26	37.84	42	2555	665.0	
9	-122.25	37.84	52	3549	707.0	

	population	households	median_income	ocean_proximity	median_house_value	\
0	322	126	8.3252	NEAR BAY	452600	
1	2401	1138	8.3014	NEAR BAY	358500	
2	496	177	7.2574	NEAR BAY	352100	
3	558	219	5.6431	NEAR BAY	341300	
4	565	259	3.8462	NEAR BAY	342200	
5	413	193	4.0368	NEAR BAY	269700	
6	1094	514	3.6591	NEAR BAY	299200	
7	1157	647	3.1200	NEAR BAY	241400	
8	1206	595	2.0804	NEAR BAY	226700	
9	1551	714	3.6912	NEAR BAY	261100	

	ocean_Dist
0	3
1	3
2	3
3	3
4	3
5	3
6	3
7	3
8	3
9	3

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	

	population	households	median_income	median_house_value	ocean_Dist
0	322	126	8.3252	452600	3
1	2401	1138	8.3014	358500	3
2	496	177	7.2574	352100	3
3	558	219	5.6431	341300	3
4	565	259	3.8462	342200	3

```
[36]: # import model
from sklearn.linear_model import LinearRegression

# instantiate
linreg = LinearRegression()

# fit the model to the training data
linreg.fit(X_train, Y_train)

#print the intercept and the coeff
print (linreg.intercept_)
print (linreg.coef_)
```

```
[-2253987.60135963]
[[-2.65022024e+04 -2.49155605e+04  1.10597481e+03 -5.87232467e+00
  1.05721352e+02 -3.73269694e+01  4.23132106e+01  3.92696824e+04
  7.26435303e+03 -4.83404123e+04  1.21359840e+05 -2.78591058e+04
 -2.58843992e+04]]
```

```
[11]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

#display (X_train[:5])
#print ()

X_train.total_bedrooms = scaler.fit_transform(X_train[['total_bedrooms']])
X_train.total_rooms = scaler.fit_transform(X_train[['total_rooms']]) # x is_
↪ column then (x- mean(x))/sd(x)
#display (X_train[:5])
X_train.population = scaler.fit_transform(X_train[['population']])
X_train.households = scaler.fit_transform(X_train[['households']])
X_train.longitude = scaler.fit_transform(X_train[['longitude']])
X_train.latitude = scaler.fit_transform(X_train[['latitude']])
```

```

X_train.housing_median_age= scaler.
↳fit_transform(X_train[['housing_median_age']])
X_train.median_income= scaler.fit_transform(X_train[['median_income']])

X_test.total_bedrooms = scaler.fit_transform(X_test[['total_bedrooms']])
X_test.total_rooms = scaler.fit_transform(X_test[['total_rooms']]) # x is_
↳column then (x- mean(x))/sd(x)
#display (X_train[:5])
X_test.population = scaler.fit_transform(X_test[['population']])
X_test.households = scaler.fit_transform(X_test[['households']])
X_test.longitude = scaler.fit_transform(X_test[['longitude']])
X_test.latitude = scaler.fit_transform(X_test[['latitude']])
X_test.housing_median_age= scaler.fit_transform(X_test[['housing_median_age']])
X_test.median_income= scaler.fit_transform(X_test[['median_income']])

Y_train.median_house_value=scaler.fit_transform(Y_train[['median_house_value']])
Y_test.median_house_value=scaler.fit_transform(Y_test[['median_house_value']])

display (X_train[:5])
display (Y_train[:5])

```

/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
```

/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
```

/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
```

```
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
```

```
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[name] = value
```

```

/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

self[name] = value
/usr/local/lib/python3.7/site-packages/pandas/core/generic.py:5303:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

self[name] = value

longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
7061      0.780934 -0.805682           0.509357    -0.113242    -0.337870
14689     1.245270 -1.339473          -0.679873    -0.213566    -0.013884
17323    -0.277552 -0.496645          -0.362745    -0.482639    -0.614210
10056    -0.706938  1.690024          -1.155565    -0.848339    -0.926284
15750    -1.430902  0.992350           1.857152     0.251071     0.400626

population  households  median_income  ocean_Dist  ocean_dist_1  \
7061      -0.184117   -0.243508      0.133506         0         0
14689     -0.376191   -0.013267     -0.532218         4         0
17323     -0.611240   -0.565322      0.170990         4         0
10056     -0.987495   -0.949929     -0.402916         1         1
15750      0.086015    0.426285     -0.299285         3         0

ocean_dist_2  ocean_dist_3  ocean_dist_4
7061          0           0           0
14689          0           0           1
17323          0           0           1
10056          0           0           0
15750          0           1           0

median_house_value
7061             -0.113387

```

14689	-0.321603
17323	0.456831
10056	-0.611895
15750	2.532087

```
[11]: display (X_train[:5])
      display (Y_train[:5])
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
7061	0.780934	-0.805682	0.509357	-0.113242	-0.337870	
14689	1.245270	-1.339473	-0.679873	-0.213566	-0.013884	
17323	-0.277552	-0.496645	-0.362745	-0.482639	-0.614210	
10056	-0.706938	1.690024	-1.155565	-0.848339	-0.926284	
15750	-1.430902	0.992350	1.857152	0.251071	0.400626	

	population	households	median_income	ocean_Dist	ocean_dist_1	\
7061	-0.184117	-0.243508	0.133506	0	0	
14689	-0.376191	-0.013267	-0.532218	4	0	
17323	-0.611240	-0.565322	0.170990	4	0	
10056	-0.987495	-0.949929	-0.402916	1	1	
15750	0.086015	0.426285	-0.299285	3	0	

	ocean_dist_2	ocean_dist_3	ocean_dist_4
7061	0	0	0
14689	0	0	1
17323	0	0	1
10056	0	0	0
15750	0	1	0

	median_house_value
7061	-0.113387
14689	-0.321603
17323	0.456831
10056	-0.611895
15750	2.532087

```
[12]: display (X_test[:5])
      display (Y_test[:5])
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
20046	0.262460	0.221006	-0.302678	-0.498771	0.010928	
3024	0.037969	-0.209747	0.096730	0.147762	0.010928	
15663	-1.448662	1.035689	1.854125	0.546562	0.010928	
20484	0.407132	-0.612406	-0.941731	0.196319	0.010928	
9814	-1.194239	0.483203	0.416256	-0.118405	0.010928	

	population	households	median_income	ocean_Dist	ocean_dist_1	\
20046	-0.024915	-0.358030	-1.152099	1	1	
3024	0.130364	0.230115	-0.701791	1	1	
15663	-0.098516	1.220811	-0.199201	3	0	
20484	0.256024	-0.002529	0.996622	0	0	
9814	-0.320216	-0.177665	-0.069475	4	0	

	ocean_dist_2	ocean_dist_3	ocean_dist_4
20046	0	0	0
3024	0	0	0
15663	0	1	0
20484	0	0	0
9814	0	0	1

	median_house_value
20046	-1.387817
3024	-1.404401
15663	2.560130
20484	0.103898
9814	0.622376

```
[12]: # import model
from sklearn.linear_model import LinearRegression

# instantiate
linreg = LinearRegression()

# fit the model to the training data
linreg.fit(X_train, Y_train)

#print the intercept and the coeff
print (linreg.intercept_)
print (linreg.coef_)
```

```
[0.11545353]
[[-0.45859533 -0.4597284  0.12052227 -0.10973858  0.38341869 -0.36770133
  0.13972452  0.64626901  0.06276157 -0.41764494  1.04850829 -0.24069332
 -0.22363252]]
```

```
[16]: # predictions on the testing data
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
Y_pred = linreg.predict(X_test)

print ("Rsqd = ", r2_score(Y_pred, Y_test))
```

```

#print("RMSE = ", np.sqrt(mean_squared_error(Y_pred, Y_test)))

#from sklearn.metrics import r2_score
#print ("Rsqd = ", r2_score(y_pred, y_test))

from sklearn.metrics import mean_squared_error
import numpy as np
print("RMSE = ", np.sqrt(mean_squared_error(Y_pred, Y_test))

```

```

Rsqd = 0.4460164533413796
RMSE = 0.6003199071326119

```

```
[18]: print(Y_pred)
```

```

[[-1.22438041]
 [-0.63011131]
 [ 0.49699754]
 ...
 [ 0.65545862]
 [-0.8198401 ]
 [ 0.09491845]]

```

```
[ ]:
```